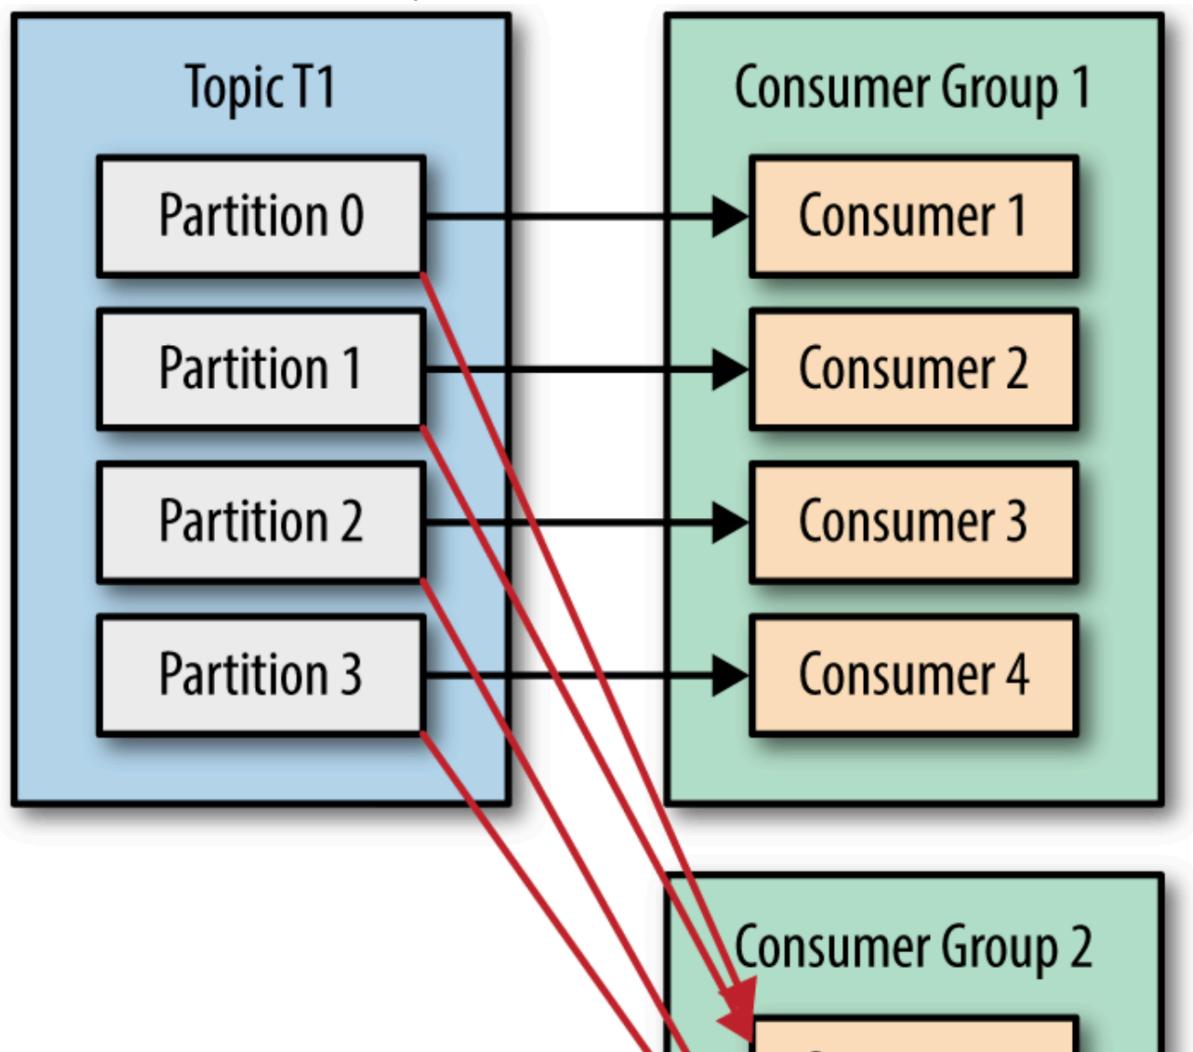


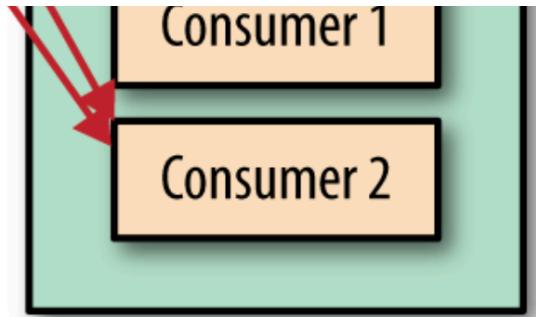
Kafka

Wednesday, 15 May 2024 8:34 AM

- Kafka organizes messages into Topics
- Each topic has a set of partitions
- Each partition is a replicated log of messages called a replicated log, with a reference to the previous offset
- Each partition is replicated 3 times
- Each replica is stored on a separate broker
- The leader handles all read and write processes
- Zookeeper manages the replication of partitions
- If you have 2 consumers with the same group ID and 12 partitions, each consumer will read 6 partitions each



each



Within same group: **NO**

Pache Matching : It is some thing producer produces more messages that cannot be handled by consumer then we can place messages in que so that consumer can process one by one
P2P: Here one producer is publishing the message then it will be read by one consumer then second consumer subscribed to same topic then only one consumer can read this message.

POP/SUB: Here one message can be published to multiple Q's through exchange
 If partition is not defined then it will follow round robin logic to place a message in partition
 Example Topic has 3 partition then first message will go to partition 0 then 2 and 3 . 4th message will come to partition 0

Consumer group is needed is because if we configure multiple consumer inside consumer group then in any case consumer one goes down then Zookeeper will redirect to other consumer to read data from where primary consumer is left over .

We can configure the topic to retry mechanism. If that does not work after n number of retries that is been configured then it will put the message dead letter topics
 Kafka uses pull back mechanism but rabbit mQ uses push back mechanism

Offset explorer UI used in Kafka to view the lag information and partition information where data is stored

Branching Strategy

Kafka Streams

Below Link for

<https://rathod-ajay.medium.com/kafka-interview-question-for-experienced-developer-in-2024-01b95d126cdd>

Stream Processing Characteristics

- 1) Unbounded data sets : While processing the data sets , It is not known how many data sets are present and expected to come
- 2) One record at a time . Each record is analyzed , transformed

d by

message

roup
co

y

e

a

- ↳ One record at a time . Each record is analyzed , transformed
- 3) Real time computation : Recorded are processed inside generics and pushed to next step
- 4) In real time
- 5) Low latency
- 5) Parallel processing

<https://developer.confluent.io/courses/spring/spring-boot-with-apache-kafka-hands-on/>

Apache kafka is a highly scalable and distributed platform for creating and processing the data in real time

Open sourced in 2011 and developed by LinkedIn

Kafka has 3 components

- 1) Kafka Connect
- 2) Kafka streams - Real time stream processing
- 3) KSQL- DB
- 4) Kafka Broker
- 5) Kafka client- For producer and consumer

Topic is data stream and we can compare to database table

Kafka cluster is nothing but a group of computers that contain the kafka servers. Whenever a topic is created, it is partitioned into multiple partitions, and each partition is replicated across multiple servers in the cluster.

Kafka consumer group is nothing but a group of consumers but max consumer can be less than the number of partitions . If consumer is more than the partitions , it will not accept if consumer is more than the partition . This is due to double rating of records

Kafka connect : is between data source and kafka cluster on source and destination , Kafka connect is a Java library that allows you to connect your data source to your target system database . As being the developer we don't need to write single line of code

The source connector is called Source connector and destination is called as Sink connector

Kafka Connect is called a cluster . Each unit in the cluster is called as worker .

In Kafka connect(Source and Sink) both we can apply the transformation logic like modify , filter , etc

Kafka connect will self manage following activities. When any worker goes down with in the cluster , it will be automatically taken care by the other workers . In this way load will be balanced

- 1. Reliability
- 2. High Availability
- 3. Scalability
- 4. Load Balancing

Inside the kafka connect we need to configure the JDBC connection and define the task (consumer and producer) . Once the task is assigned to the kafka connect cluster once processing is complete it task will provide the update to the database

stage

streams in real time

partition is created kafka server will distribute the data of

than or equal to number partition of the TOPIC . Kafka will
s or not to process the duplicate the records.

kafka connect will read data from kafka cluster and post in to
table

actor .This is uses Kafka connect frame work

ying the records

Kafka connect cluster the task are reassigned to other

configure the the tables) then task will be distributed across
worker

KSQL : It is SQL interface to Kafka Stream . It has below 2 modes

- 1)Interactive mode- command line interface or Web interface we can submit the query .
- 2) Headless Mode: Here we can submit the KSQL file which are executed by KSQL server .

KSQL comes with 3 components

- 3) KSQL Engine
- 4) REST Interface
- 5) KSQL Client(CLI/UI)

Combination of KSQL Engine and Rest interface form the KSQL Server . It can be deployed

Using KSQL you can group , apply filter , join topics & aggregate the topics

Sink the result of your query to another topic

Kafka comes with 3 different flavor

- 1) Open Source - Apache kafka
- 2) Commercial Distribution - confluent .io - Used in prod , it is also managed service
- 3) Managed service -confluent, amazon, aiven.io - U don't need to download the software provider

From confluent.io website and hit on download button

To start the kafka zookeeper

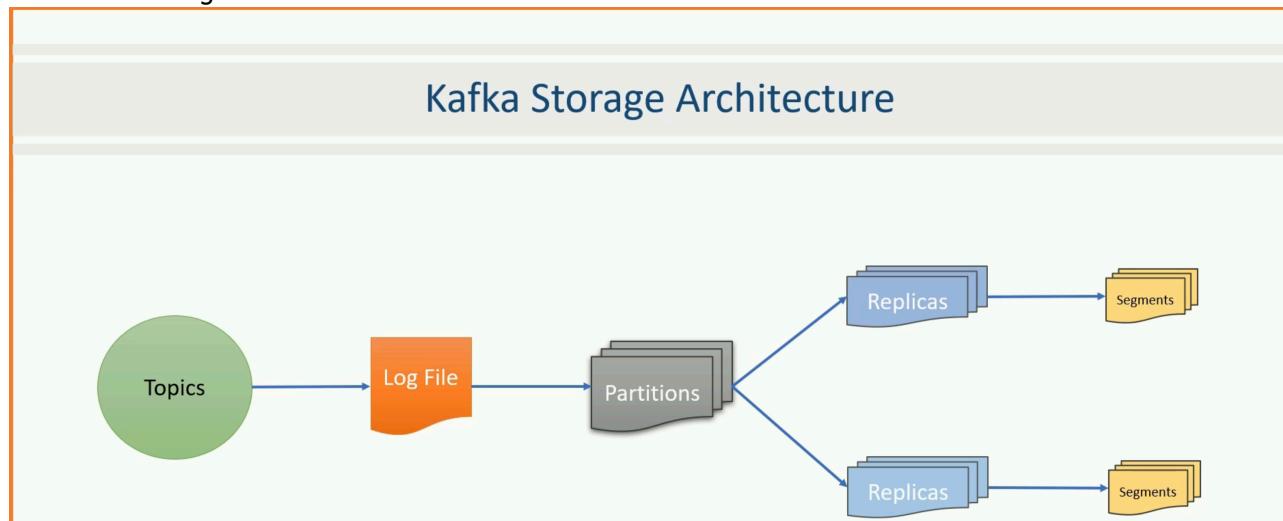
```
sh zookeeper-server-start /etc/kafka/zookeeper.properties
```

Below are the properties to start multiple instance of kafka (kafka cluster)

Under this path /Downloads/confluent-7.6.1/etc/kafka you can find server. properties you

- Broker-id =0 by default you can change it to 1 , 2 and so on but in real time you must have separate machine
- Listener port you need to change by default it is #listeners=PLAINTEXT://:9092 uncontrolled port
- Logfile name you need to change log file log.dirs=/tmp/kafka-logs you can change it

Kafka Storage Architecture



Good for development and test env

Good for prod env

in interactive mode or headless mode

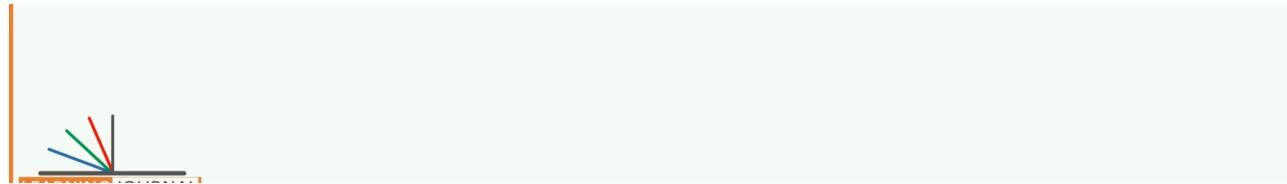
and manage the cluster . All these will be managed by cloud

you need to make following changes

you not need to do it because each server is started on

environment and change respectively

change to log.dirs=/tmp/kafka-logs1 etc



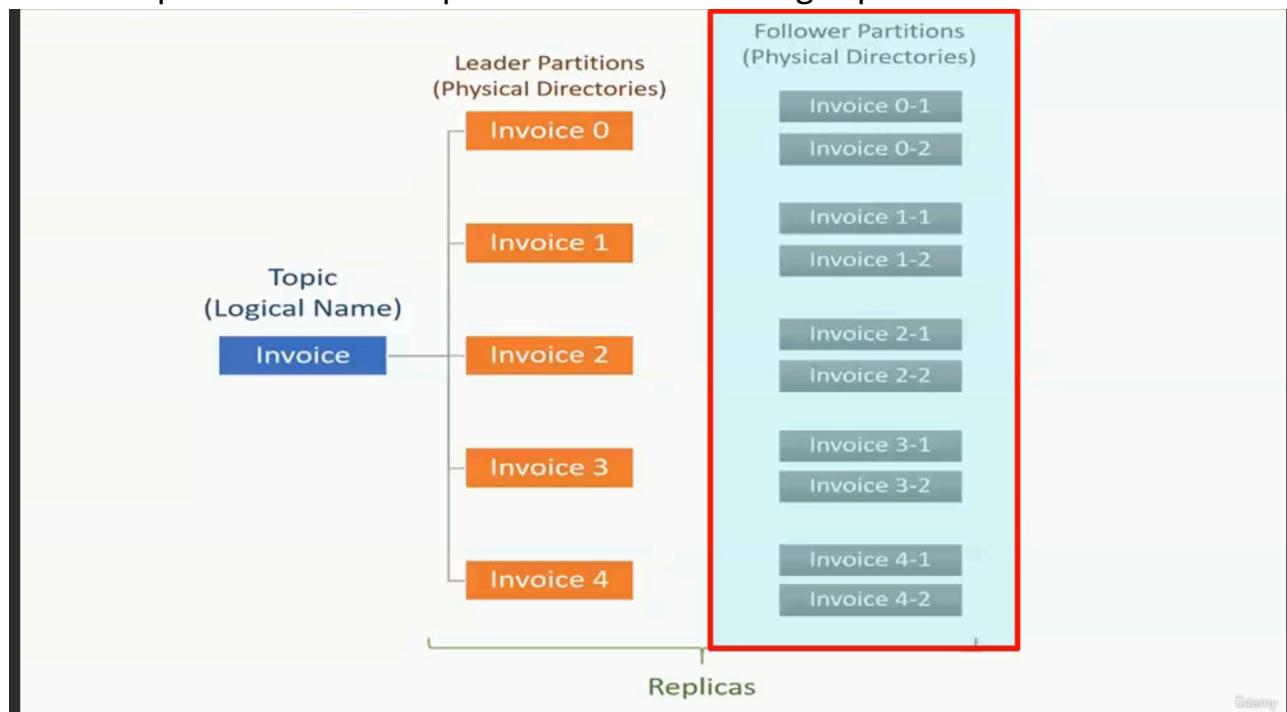
Topic will be divided into partitions. Example: million messages are not able to keep at single partition.

Replication factor is nothing but how many copies you need to keep at each partition. For 5 partitions and 3 replicas, then 15 directories should be created. In this example, 5 directories are shown.

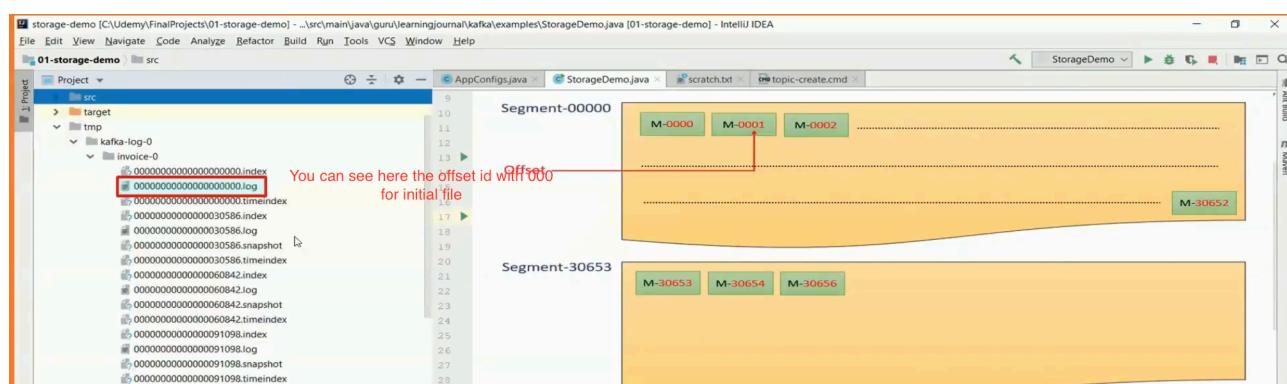
We can categorize the topic partition replicas into 2 categories:

Leader partition - each partition created during topic creation is leader.

Follower partition - each replica set created during topic creation is follower.



Offset message is of 64-bit integer giving the unique ID for each message. In case of reading, it starts from the next offset ID from the last segment. You can see in below screen shot initial .log file is started.



single place and divided and placed at multiple place called

for example while creating the topic you have specified the 5
partitions and each broker will be created in case if you started 3 Kafka

partitions the max size then new segment will be created with
the previous segment will be closed with zero and next log file with some numbers

```

leader-epoch-checkpoint
> invoice-1
> invoice-2
> invoice-3
Run  zookeeper-start  0-kafka-server-start  1-kafka-server-start  2-kafka-server-start  StorageDemo  main0
29
30
-->
StorageDemo  main0

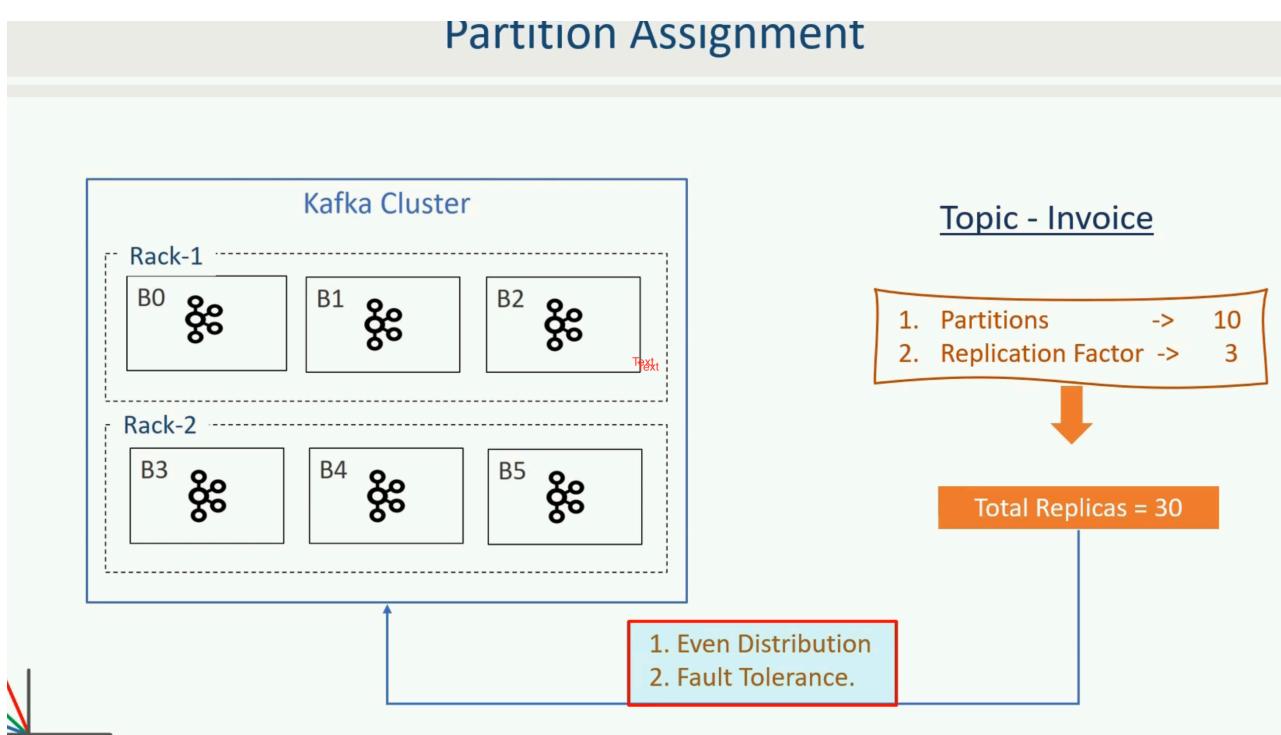
```

In the above screen shot .index will main the off set index of messages and .timeindex will maintain the time stamp of message from certain time

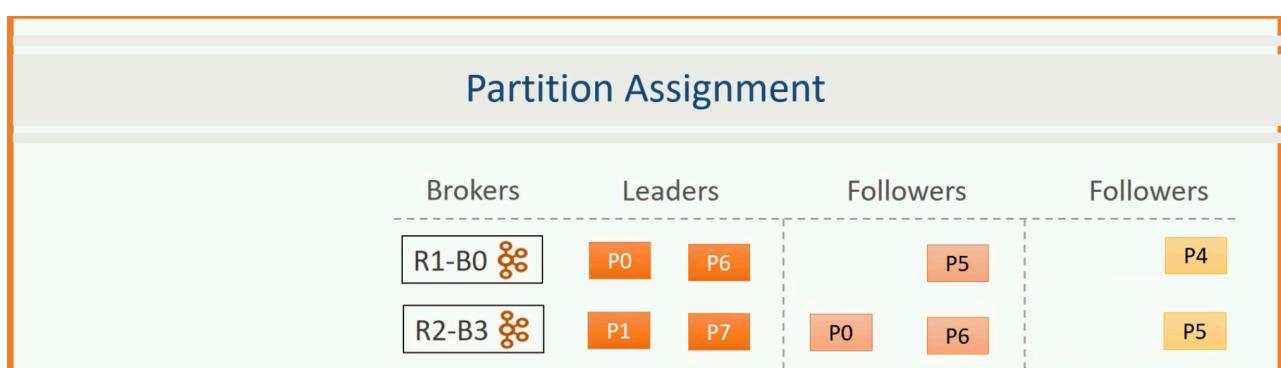
Kafka Cluster is nothing but group of brokers created and work together to perform the tasks . We can create 100 plus broker in single cluster .

Controller is managing the brokers in the cluster , If any broker dies or left cluster it will reassign the leader to another broker to become the lead broker

Partition assignment is shown as per the below diagram , For example if the Kafka Cluster has 10 partitions and replication factor of 3 . Total replicas will be 30 . This will ensure scalability and copy of the request will be saved in different machines .



This is how Brokers, leader and followers are distributed in the cluster . First Brokers are assigned to racks . Then leaders are assigned to each broker . Finally followers will be assigned in the fashion .



ll maintain the time . In case consumer wants to read

ask

eassign the work to other broker . So that other broker

is created with 6 broker and placed on 2 racks .

partition is 30 will be distributed even as possible to make
one to make sure to achieve the fault tolerance

arranged in order like one broker from one rack .

of leaving first broker



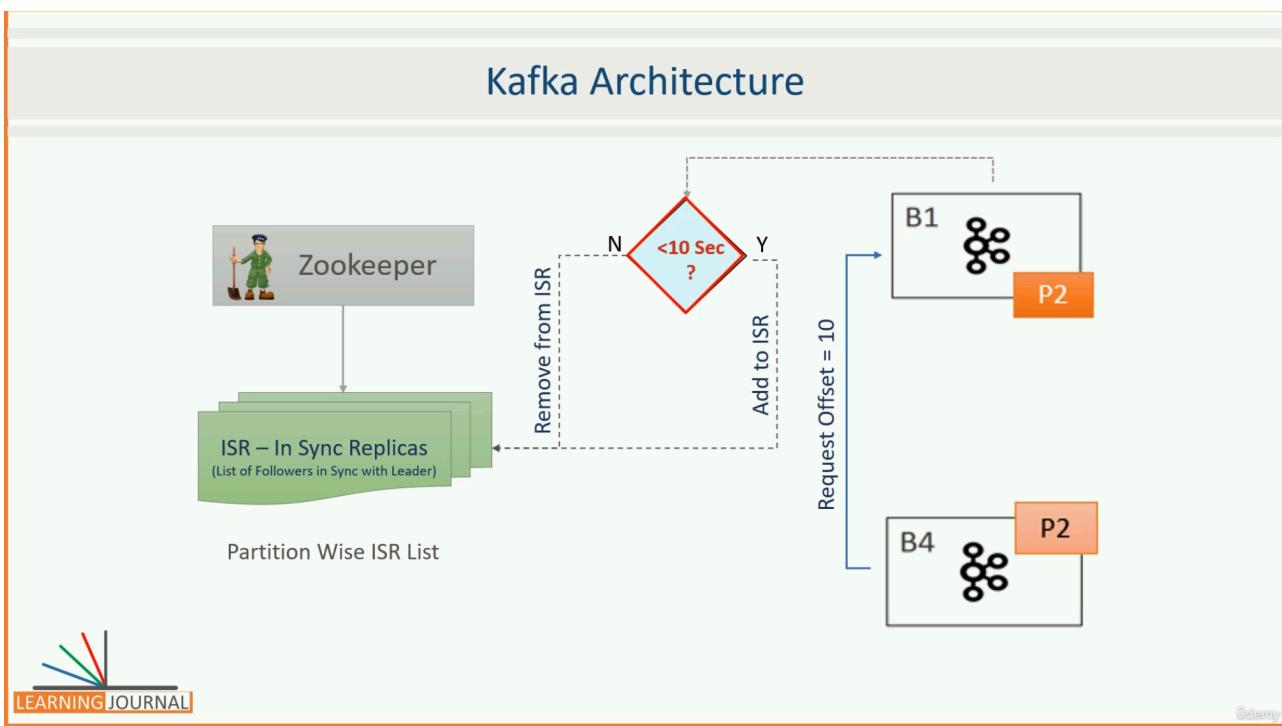
How does Producer make a request to Topic.

Producer will query for broker in Kafka cluster and it will respond back with meta data wi

Producer will decide which partition it needs to send data . On receiving end also it will re

Followers will not accept message from producer or consumer . Followers will request the

ISR List-In Sync Replica : ISR refers to the replicas of a partition that are “in sync” with the leader. As soon as follower request for message ,it will send from lead broker . First time follower offset will be from 10 onwards. So that leader will understand the message from 0 to 9 was per In this scenario there some time needed to sync between leader and broker that time by configuration . In sync replicas contain the followers who are sink with leader with in config



Committed vs uncommitted message: The message which is successfully sent to follower is called committed message.

Minimum In Sync replicas : It is managed by kafka incase Kafka brokers went down for some reason. It will exists only in leader. It is a risky situation in order to avoid this we need to set min-in-sync-replicas. The negative thing is if we set the min-insync-replicas to 2 then until 2 replicas commit the message it will be in read only mode.

with all leaders .

read from Leader Broker.

the message and leader will send a message to stay in sink

with the leader.

Broker will make request from off set 0 to 9 . Then next request

persist .

default is 10 sec . You can increase or decrease using Kafka
configured seconds .(This scenario it is 10 sec).

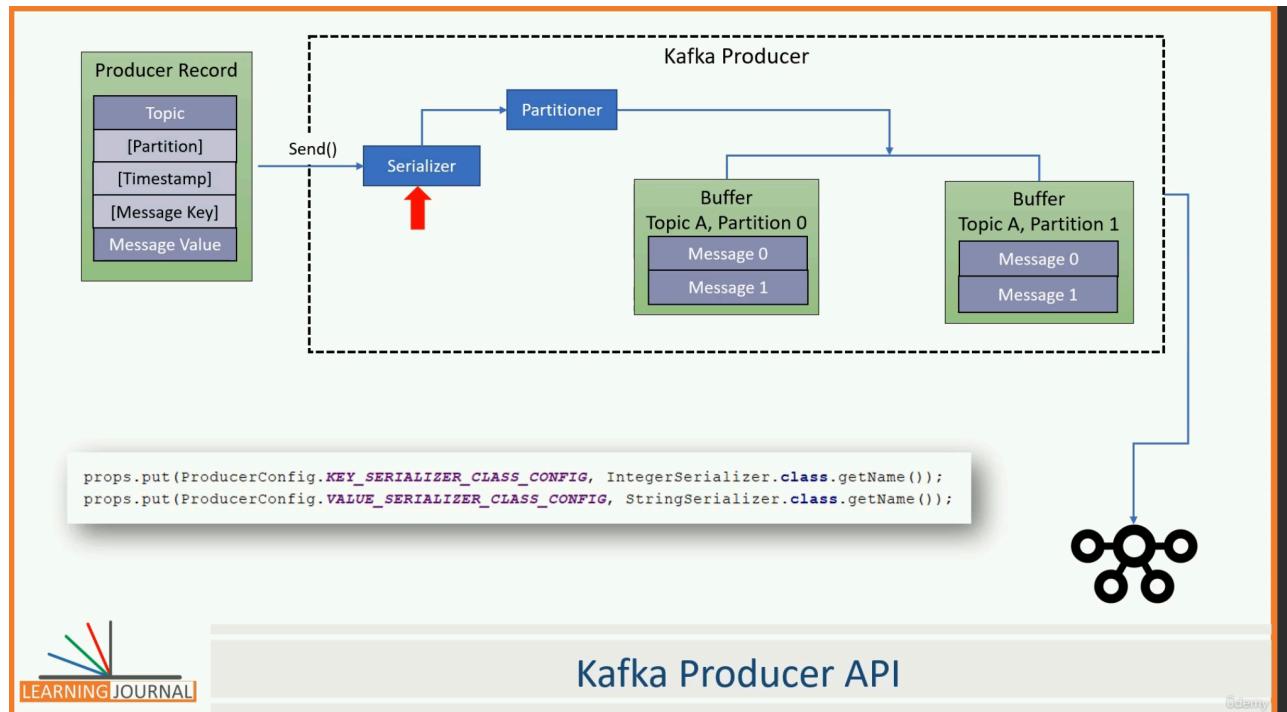
is called committed message. Which is not in follower is

some reason and left with only Leader broker . Then our data
sync-replicas : 2.

the message . Topic will not allow to write message and it is in

Kafka Producer API

In producer API there are 5 parameters , 2 are mandatory (topic name , partition) .Below is the producer architecture . Before the message needs to be serialized .For complex serialization you can use avro , json etc



Kafka broker has multiple partition . In order to send the message to specific partition we need to specify the partition parameter while sending the request .

Second option is we need to specify the key in the request then Kafka will automatically place the message in the correct partition .

By default kafka has default partitioner (Hash Key Partitioning , Round Robin Partitioning) . In this partitioner Kafka uses the key and find the hash code associated with key and ideally messages with same key goes to same partition . This algorithm takes total number of partitions into account and if we increase the partition then default partitioner will start giving different partition . If we want to have more partitions initially create enough partition or 25 % more . The disadvantage is if we want to change the partition then we need to re publish the messages.

In case if the key is not specified then default partitioner will use round robin partitioning and each message will be sent to each broker

Once message is reached the Buffer , I/O thread will transfer the message to the broker

key and message) optionals are (key, timestamp ,
needs to be send to broker it needs to be

specific partition you can specify the partition

will use the key to find the partition where

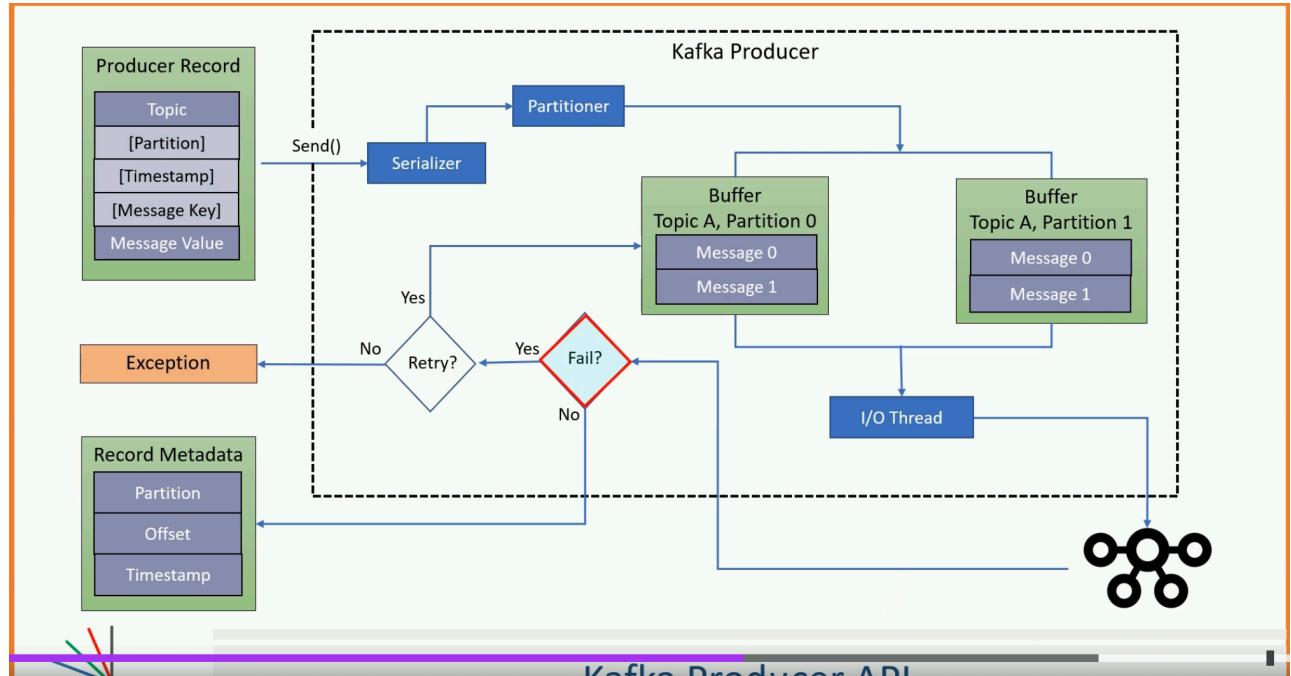
Round robin Partitioning) .

Identify partition . Hashing ensures same key
of partition as another argument . If you
put out put. Means while creating the partition
you increase the partition drastically then you

Round robin process to distribute the messages to

message to topic . If buffer memory is exhausted

then send method will get I/O exception in that case we need to increase Memory in configuration class . Default memory size is 32 mb. Complete flow is below



At least once & At Most Once

At least once: Once producer send the message it will be stored in kafka success acknowledgment from broker . IO will try to resend a message but it will not reach the IO due to network error . In this case producer network error . In this scenario Kafka does not have mechanism to identify same message to log file . This implementation is known as atleast once.

Atmost once means can be achieved by setting the retry to zero . In this case we have duplicate.

We need to set this one as part of properties configuration in producer Configuration:

java

Copy code

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("key.serializer", "org.apache.kafka.common.serialization.StringSeriali
props.put("value.serializer", "org.apache.kafka.common.serialization.StringSeriali
props.put("acks", "0"); // No acknowledgment
props.put("retries", "0"); // No retries
props.put("delivery.timeout.ms", "30000"); // Maximum wait time for acknowledgm
KafkaProducer<String, String> producer = new KafkaProducer<>(props);
```

use the memory of Producer using buffer.

ka log file and incase IO thread did not get the
e this time broker sends success message to IO
r will try to resend the same message due to
entify the duplicate entry it and it will store
ce
this scenario you many loose messages but never

er API . Below is the example.

Consumer Configuration:

java

 Copy code

```
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("group.id", "my-group");
props.put("key.deserializer", "org.apache.kafka.common.serialization.StringDeser");
props.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeser");
props.put("enable.auto.commit", "true"); // Auto commit offsets
props.put("auto.commit.interval.ms", "5000"); // Commit interval
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);
```



Exactly Once : By enabling the property called `enable.idempotence = true` in consumer configuration. This is achieved by implementing exactly once

Initially when producer is handshake with broker. Producer API will assign unique producer ID for each producer.

Next step is producer API will start assign the unique message sequence number. When producer send message it will uniquely identify the message using producer ID and sequence number.

POS(Point of Sale) simulator in Kafka: It basically helps us to send millions of messages to Kafka. It takes arguments

- Topic name: Name of the topic where we need to send message
- Number of Producer thread: How many parallel thread needs to be run
- Produce speed: Where we can specify the time interval between two messages

Note: Size of the message we cannot specify that is the draw back

Need to do practical practise

Kafka Consumer Example architecture

We need to define consumer group that contains the multiple consumers. If we have 5 consumer in consumer group and 10 partitions in same topic then kafka will take care of scalability and fault tolerance. If we have 5 consumer and 10 partitions in topic then each consumer will consume 2 partitions per second for 10 topic then target consumer group has 5 consumer to consume 10 topics. In the scenario like 10 consumer and 10 partitions then kafka will assign the current consumer topic to different consumer.

Kafka Consumer Application

true. If we enable this kafka take care of

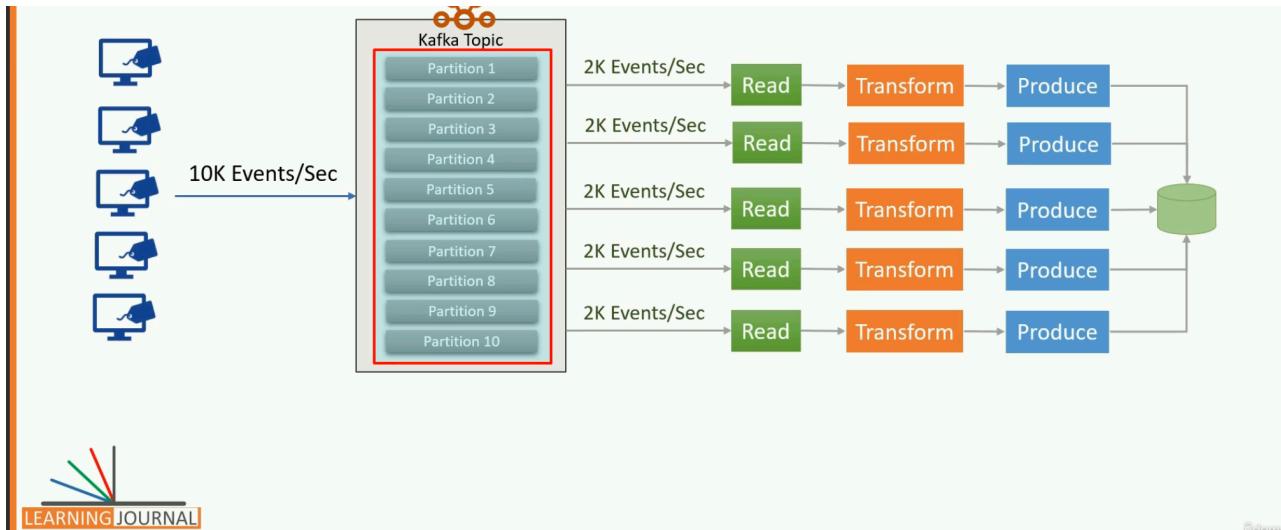
k for ID and broker will assign the dynamic

ce number . It starts from zero. When IO thred
and sequence number

ons of messages in parallel. It takes 3 major

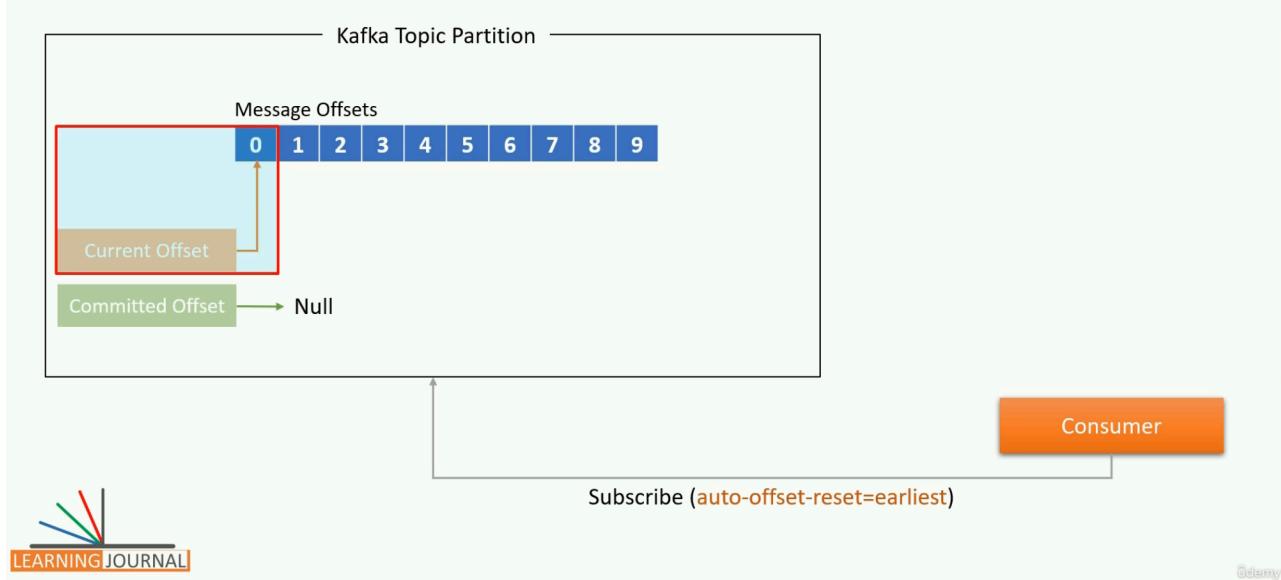
be enabled to send message
the messages.

mers. When consumer group is subscribed to
In the below picture if 10k events are produced
hen 2k messages will be distributed for each
n . If any consumer goes down then Kafka will re



Kafka maintains the unique records using the offset ID & also current partition. Initially the current offset will be null . It means consumer no. You can set the property `auto-offset-reset=earliest`. It means current offset means will send only upcoming messages after the subscription.

Kafka Consumer Application



Once consumer start reading the messages the offset will be advanced . offset will be set to 4 . In case of consumer restarts the offset will reset again that is duplicate. To avoid that situation kafka maintains the consumer offset in log. If consumer offset is not available then consumer will poll down then other consumer can read data by using commit offset number. Note: The committed offset will be updated at broker only after consumer has read data base or different mS while going to poll(read) next message .

offset and committed offset with in the
needs to start reading from offset number 0 .
offset is set to zero. The default value latest that

d . Example if 3 messages are consumed then
set and it will result in getting same message
committed offset . It will help if the consumer goes
number as it is securely stored with broker.
sumer is successfully read and update the target

Learn : <https://www.linkedin.com/in/prashant-kumar-pandey/>

